

Statement of Purpose

Implementing a CI/CD Pipeline Using Azure DevOps

By Arun Nambiar and Hemanth Kumar Kotha

Statement of Intent:

This is written with intent to giving a thorough guide around how to implement a Continuous Integration and Continuous Deployment (CI/CD) pipeline utilizing Azure DevOps. CI/CD pipelines play a crucial role in ensuring rapid and reliable software delivery. We are excited to contribute to the COE and assisting members in creating seamless development and deployment processes within the software development life-cycle.

Background and Motivation:

In today's fast-paced software development landscape, the adoption of CI/CD methodologies has become imperative to maintain competitiveness and deliver value to end-users. Recognizing the significance of automating build, testing, and deployment processes, We are motivated to implement a robust CI/CD pipeline to facilitate the swift delivery of high-quality software. Azure DevOps offers a comprehensive suite of tools to achieve this goal effectively.

Objectives:

- 1) Automated Build and Testing: We aim to leverage Azure DevOps to automate the build and testing stages of the software development process. This will involve configuring the pipeline to trigger builds automatically whenever changes are pushed to the repository. By integrating unit tests, integration tests, and other relevant testing frameworks, We intend to ensure that code quality is maintained throughout the development life-cycle.
- 2) Continuous Deployment: Our objective is to implement a seamless deployment process using Azure DevOps, enabling automated deployment of applications to various environments, such as development, staging, and production. This will involve setting up deployment slots, environment-specific configurations, and approval gates to ensure smooth transitions between stages.
- 3) Monitoring and Reporting: We plan to incorporate monitoring and reporting mechanisms within the pipeline using Azure DevOps tools. This will enable real-time visibility into the pipeline's health, test results, and deployment status. By integrating monitoring tools, I aim to promptly identify and address any issues that may arise during the deployment process.

Approach:

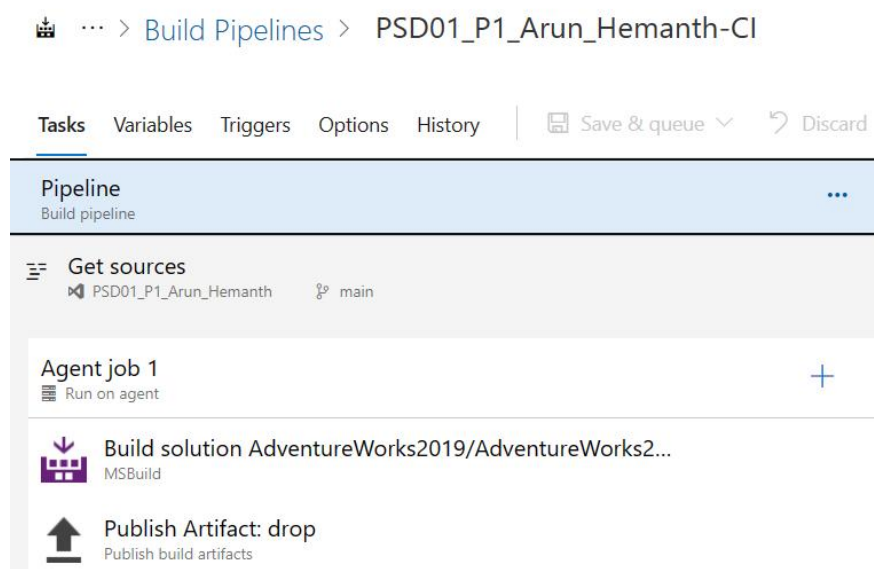
(Brief description of the process, for detailed version kindly refer to the Proof of concept at the end of this document)

1) Pipeline Configurations: We will start by creating a new project in Azure DevOps and setting up the necessary repository for version control. We will configure build pipelines using Azure Pipelines, defining build steps, and incorporating code analysis tools for quality checks.

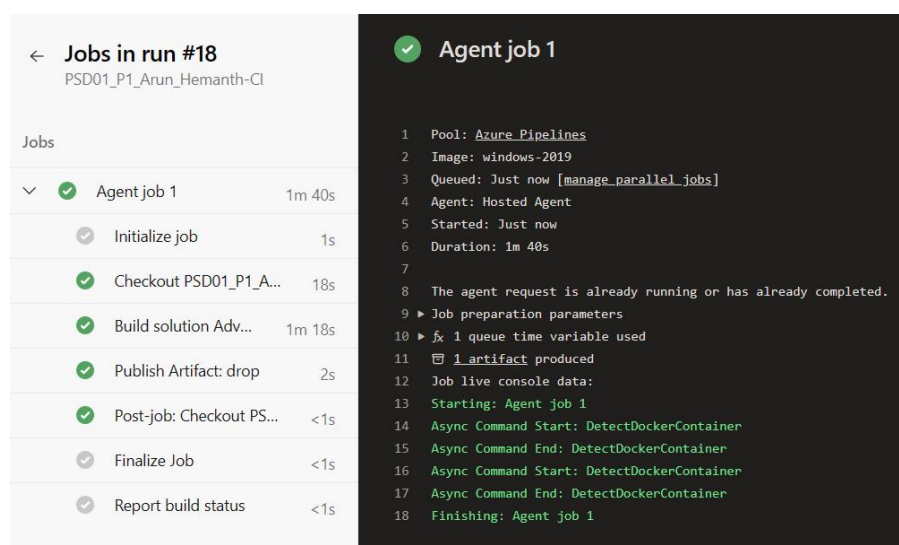
Jobs required for the agent:

1.1) Build solution task

1.2) Publish Artifact : Drop



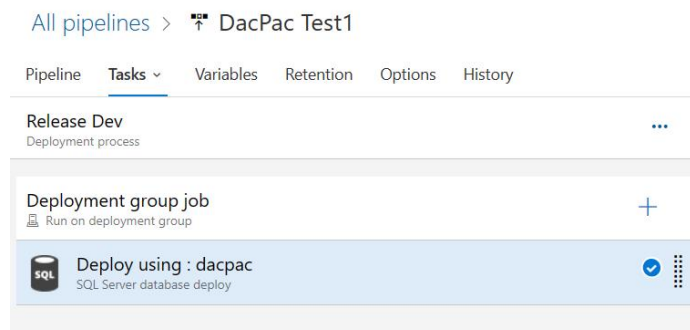
2) Testing Integration: For testing you can add files to your GIT and see if that triggers the build pipeline, if it does, we can check for the artifact drop. If the build pipeline runs seamlessly then we will have an artifact drop with the new solution file. Using this we aim to ensure the reliability and stability of the software.



3) Deployment Setups: We will set up deployment/release pipelines, defining deployment stages and environment configurations. Azure DevOps provides integration with Azure services, allowing us to deploy applications to Azure App Services, Kubernetes clusters, or other relevant platforms.

Jobs required for the agent:

3.1) Deploy using DACPAC



The Task we used has the following properties:

SQL Server database deploy ⓘ [View YAML](#) [Remove](#)

Task version 0.* ▼

Display name *
Deploy using : dacpac

Deploy SQL Using * ⓘ
Sql Dacpac ▼

DACPAC File * ⓘ
\$(System.DefaultWorkingDirectory)/_EDM Impact-
CI/drop/AdventureWorks2019.dacpac ...

Specify SQL Using * ⓘ
Connection String ▼

Connection String * ⓘ
\$(ConnectionString)

Publish Profile ⓘ
\$(System.DefaultWorkingDirectory)/_EDM Impact-
CI/drop/PublishProfiles/AdventureWorks2019.publish.xml

Additional Arguments ⓘ
/p:DropObjectsNotInSource=True /p:BlockOnPossibleDataLoss=False

Control Options ▼

Output Variables ▼

4) Monitoring Implementations: Using Release Monitor, we will incorporate monitoring and logging capabilities into the pipeline. This will enable the tracking of application performance, error rates, and other key metrics.

DacPac Test1

Releases
Deployments
Analytics

Releases	Created	Stages
<div>HK</div> Release-89 77 main	7/5/2023, 3:51:55 PM	<div>✓ Release D...</div>
<div>HK</div> Release-88 76 main	7/3/2023, 5:17:38 PM	<div>✓ Release D...</div>
<div>HK</div> Release-87 75 main	6/14/2023, 8:59:02 PM	<div>✗ Release D...</div>

DacPac Test1 > Release-88 > Release Dev ✓ Succeeded

← Pipeline Tasks Variables **Logs** Tests
Deploy Cancel Refresh Download all logs

Deployment process
Succeeded

✓ Deployment group job
Succeeded

Deployment group job
Deployment group: LOCAL

1

0

SUCCEEDED

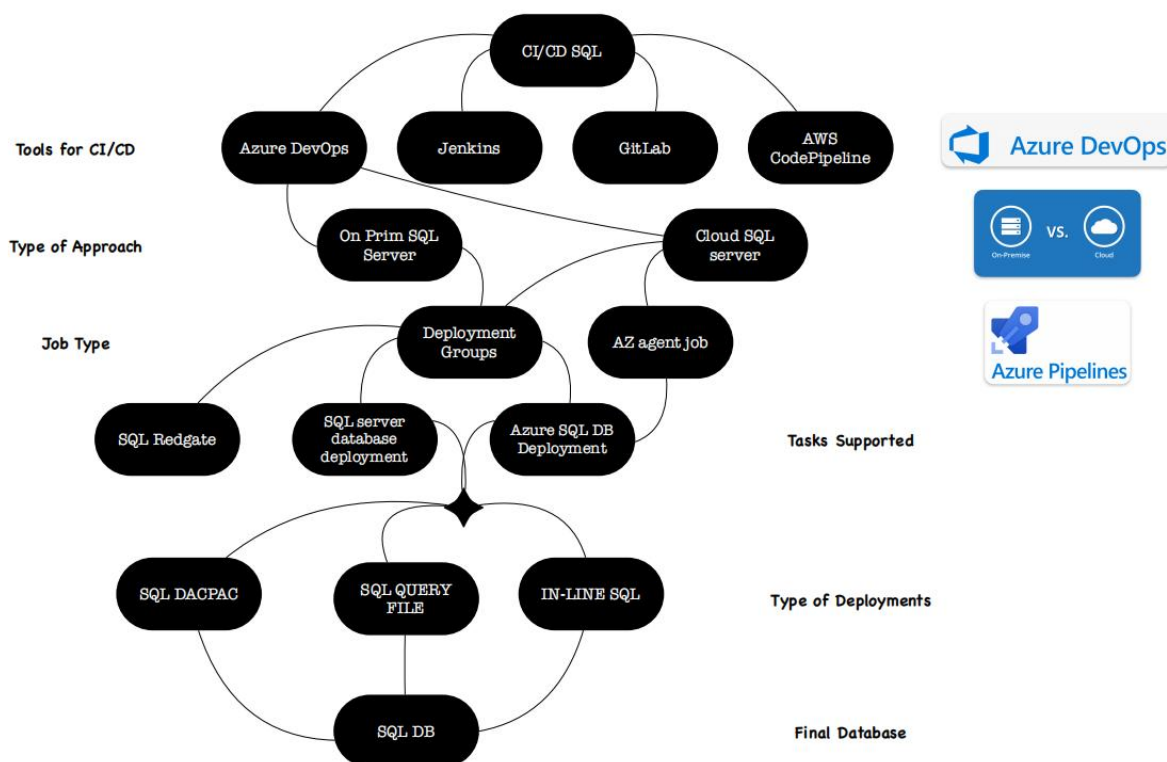
FAILED

100% PASSED

Conclusion:

In conclusion, we are enthusiastic about the opportunity to implement a robust CI/CD pipeline using Azure DevOps. This endeavor aligns perfectly with our career aspirations to excel in software development practices and contribute meaningfully to the advancement of efficient software delivery processes. We are committed to acquiring the necessary skills and knowledge to successfully design, implement, and maintain a CI/CD pipeline that enhances software quality, accelerates deployment cycles, and fosters collaboration within development teams.

On Premesis V/S Cloud



On-Premises CI/CD Pipelines:

Advantages:

1. **Control:** With on-premises pipelines, you have complete control over your infrastructure, security, and data. This is particularly important for organizations with strict compliance requirements.
2. **Customization:** You can tailor the environment to your specific needs, accommodating unique security measures, network configurations, and tool integrations.
3. **Data Location:** Data stays within your organization's network, which might be a legal or regulatory requirement for certain industries or regions.

Considerations:

1. **Maintenance:** On-premises infrastructure requires ongoing maintenance, including updates, patches, and hardware upgrades. This can be resource-intensive.
2. **Scalability:** Scaling up resources might be more complex and time-consuming compared to cloud solutions.
3. **Initial Investment:** Setting up on-premises infrastructure can involve higher initial costs for hardware and software licenses.

Cloud-Based CI/CD Pipelines:

Advantages:

1. **Scalability:** Cloud solutions offer dynamic scalability, allowing you to scale resources up or down based on demand. This is especially useful for handling fluctuations in build and deployment workloads.
2. **Managed Services:** Cloud providers like Azure offer managed CI/CD services that handle infrastructure management, updates, and security, freeing your team from these tasks.
3. **Ease of Setup:** Cloud-based pipelines often have quicker setup times compared to configuring and maintaining on-premises infrastructure.
4. **Global Accessibility:** Cloud-based solutions can be accessed from anywhere with an internet connection, fostering collaboration among distributed teams.
5. **Pay-as-You-Go:** Cloud services are often billed on a pay-as-you-go model, potentially reducing upfront costs.

Considerations:

1. **Data Privacy:** Depending on your industry and region, data privacy regulations might impact your ability to store certain data in the cloud.
2. **Dependency on Cloud Provider:** You rely on the cloud provider's services and infrastructure. If the provider experiences downtime, your pipelines could be affected.
3. **Security Concerns:** While cloud providers invest heavily in security, some organizations may still have concerns about data breaches or unauthorized access.
4. **Costs:** While pay-as-you-go can be cost-effective, extensive usage can lead to higher monthly bills. Careful monitoring and resource management are essential.

Conclusion:

Ultimately, the choice between on-premises and cloud-based CI/CD pipelines depends on the organization's specific needs, existing infrastructure, regulatory requirements, and overall IT strategy. Many organizations today are opting for hybrid solutions that combine the benefits of both approaches to create a flexible and efficient pipeline setup.

Agents V/S Deployment Groups

In Azure DevOps, both Azure Pipelines agents and Deployment Groups are used to facilitate the deployment of applications and infrastructure. Let's explore the differences and considerations between using Azure Pipelines agents and Deployment Group jobs:

Azure Pipelines Agent:

An Azure Pipelines agent is a piece of software that runs on a target machine and is responsible for executing tasks as part of a build or release pipeline. These agents can be hosted on various platforms, including on-premises servers, virtual machines, or cloud-based resources. When you define a job in an Azure Pipeline, you specify which agent pool the job should use, and Azure DevOps will then allocate an available agent from that pool to execute the tasks defined in the job.

Advantages of Azure Pipelines Agents:

1. **Customization:** You have granular control over the environment in which tasks are executed. This allows you to configure the agent to match specific requirements of your deployment.
2. **Control:** You can control the security settings, configurations, and dependencies on the agent, making it suitable for a wide range of deployment scenarios. · ·
3. **Distributed Environments:** Useful when dealing with a complex environment that requires specific configurations for different parts of the deployment.

Deployment Group Job:

A Deployment Group job is used in Azure DevOps when you want to deploy applications to a group of target machines simultaneously. A deployment group is a logical group of machines that can be on-premises or in the cloud. When you define a Deployment Group job in a release pipeline, you specify the target deployment group, and the release tasks are executed on all machines within that group.

Advantages of Deployment Group Job:

1. **Simplicity:** Deployment Group jobs offer an abstraction layer that simplifies deploying to multiple machines as a group rather than dealing with individual agents. · ·
2. **Scalability:** You can easily add or remove machines from the deployment group, making it scalable to match your deployment needs. · ·
3. **Parallelism:** Deployment tasks can be executed in parallel across multiple machines within the deployment group, potentially improving deployment speed. · ·
4. **Visibility:** Provides better visibility into the status of deployments on different machines through the Azure DevOps interface. · ·
5. **Security:** You can manage permissions and security settings at the deployment group level.

Considerations:

1. Use Azure Pipelines agents when you need fine-grained control and customization over individual machines or environments, and when you're working in a distributed or complex environment. · ·
2. Use Deployment Group jobs when you want a higher level of abstraction, parallelism, and simplified management of deployments across a group of machines with similar purposes. · ·
3. You can mix and match both approaches within a single release pipeline, depending on your deployment needs.

Here are some considerations for handling exceptions in a SQL Server deployment project using DACPAC:

Pre-Deployment and Post-Deployment Scripts:

- ◆ DACPAC deployment projects in Visual Studio support pre-deployment and post-deployment scripts. These scripts allow you to perform custom actions before and after deploying changes. You can use these scripts to handle exceptions, data migrations, data transformations, and other tasks that are not directly supported by DACPAC.

Conditional Deployments:

- ◆ You can use SQLCMD variables in your DACPAC project to conditionally deploy or not deploy specific database objects based on certain conditions. For example, you could use variables to handle different behaviors for development, testing, and production deployments.

Object-Level Customizations:

- ◆ For specific exceptions related to individual database objects (tables, views, stored procedures, etc.), you can use the pre-deployment and post-deployment scripts to handle these scenarios. You might need to use T-SQL statements to check for certain conditions before applying changes.

Transactional Handling:

- ◆ If your deployment involves schema changes that could potentially break existing transactions, you might need to plan for handling these scenarios during deployment. You could implement a strategy to ensure that long-running transactions or active sessions are not affected by schema changes.

Backup and Rollback Strategies:

- ◆ Having a backup and rollback strategy in place is essential in case a deployment encounters exceptions or issues. You could automate database backups before deployment and implement a rollback process that restores the database to the previous state if needed.

Testing and Staging:

- ◆ Before deploying changes to a production environment, it's crucial to test deployments in a staging or testing environment. This allows you to catch exceptions and issues before they impact on the production environment.

Monitoring and Logging:

- ◆ Implement logging mechanisms that capture information about each deployment, including exceptions and errors. Monitoring tools can help identify problems early and facilitate quick resolution.

Version Control and Collaboration:

- ◆ Utilize version control to track changes to your DACPAC project. This enables collaboration and helps you understand the history of changes. Collaboration tools help teams communicate about exceptions and implement solutions collaboratively.

Detailed Steps for construction:

Proof of Concept (POC) for Creating an Instructional Manual

Objective:

The objective of this Proof of Concept is to demonstrate the feasibility and effectiveness of creating a comprehensive instructional manual for users, focusing on enhancing clarity, user-centric design, and effective communication.

Detailed Step by Step:

1) Creation and maintaining Solution files in Visual Studio Code:

In Visual Studio, you can work with SQL Server databases using SQL Server Data Tools (SSDT) by creating SQL Server projects. These projects can contain DACPACs (Database Application Packages) or SQL script files to manage your database schema and objects. Let's explore both approaches:

1.1) Using DACPACs:

Creating a Solution:

Open Visual Studio and create a new SQL Server Database Project. This project will serve as a container for your database schema, objects, and scripts. Working with DACPACs:

Generate a DACPAC:

Build your SQL Server Database Project to generate a DACPAC. This DACPAC is a binary representation of your database schema.

Extracting a DACPAC:

You can extract the schema information from a DACPAC using tools like SqlPackage.exe (a command-line utility). This allows you to view the schema as SQL script files.

Loading a DACPAC:

You can load the DACPAC back into Visual Studio or deploy it directly to a target SQL Server instance to update its schema.

1.2) Using SQL Script Files:

Creating a Solution:

Open Visual Studio and create a new SQL Server Database Project. Add SQL script files to the project, representing the database schema, objects, and other components.

Working with SQL Script Files:

Define Scripts:

Create SQL script files for each database object (tables, views, stored procedures, etc.) within the project.

Uploading Scripts:

Write or upload existing SQL script files to your SQL Server Database Project.

Building and Deploying:

You can build the project to validate the syntax and correctness of the scripts. Once validated, deploy the scripts to a target SQL Server instance to update its schema.

Advantages:

DACPAC Approach:

1. Centralized representation of the database schema in a binary format.
2. Simplifies deployment by packaging schema changes into a single DACPAC file.
3. Supports schema comparison and data-tier application management.

SQL Script Files Approach:

1. Offers more visibility into the individual SQL scripts, allowing you to manage them manually.
2. Can be more suitable for projects that require direct control over each script.

Considerations:

1. Choose the approach that aligns with your team's workflow and project requirements.
2. DACPACs abstract the schema, while script files offer more transparency and control over individual changes.
3. Regularly back up your SQL Server projects and version control repositories to prevent data loss.

2) Adding and Modifying Database objects

In a Visual Studio SQL Server Database Project, you can add new database objects and alter existing ones using SQL script files. Here's a step-by-step guide on how to add new objects and make alterations:

2.1) Adding New Database Objects:

Open Visual Studio:

Launch Visual Studio and open your SQL Server Database Project.

Add New Item:

Right-click on the project in the Solution Explorer, select "Add," and then choose the type of database object you want to add (e.g., Table, Stored Procedure, View, etc.).

Configure the Object:

A new SQL script file will be added to the project. Open this file and write the SQL script to create the new database object. For example, for a table, you might write a CREATE TABLE statement.

Build the Project:

Build the project to ensure that the new script is valid and doesn't contain syntax errors. Right-click on the project and select "Build."

Deploy the Project:

After building, you can deploy the project to a target SQL Server instance. Right-click on the project and select "Deploy." This will apply the changes to the target database.

2.2) Altering Existing Database Objects:

Open Existing Object Script:

In Solution Explorer, locate the script file for the existing database object you want to alter. Double-click to open the script.

Modify the Script:

Make the necessary changes to the SQL script to alter the object. For example, if you're altering a stored procedure, modify the script to include the changes you need.

Build the Project:

After making changes, build the project to validate the script's correctness and catch any syntax errors.

Deploy the Project:

Deploy the project again to apply the changes to the target SQL Server database.

3) Creating and Usage of Publish Profiles:

Using publish profiles in SQL Server Database Projects is a powerful way to manage variables and configuration settings when deploying your database changes in continuous integration/continuous deployment (CI/CD) pipelines. Publish profiles help you maintain different deployment configurations for different environments and manage variables specific to each environment. Here's how to create and use publish profiles effectively:

3.1) Creating Publish Profiles:

Open Visual Studio:

Launch Visual Studio and open your SQL Server Database Project.

Add Publish Profile:

In the Solution Explorer, right-click on the project, select "Publish," and then choose "New Publish Profile."

Configure Profile Settings:

In the publish profile settings, you can specify various deployment options such as target server, authentication, database name, and advanced deployment settings.

Define Deployment Variables:

Under the "Advanced" tab of the publish profile settings, you can define deployment variables specific to the environment. These can include connection strings, application settings, or any other variables your database scripts depend on.

Save the Profile:

Save the publish profile with a meaningful name that indicates the environment it's intended for (e.g., "Dev," "QA," "Production").

3.2) Using Publish Profiles in CI/CD Pipelines:

Version Control:

Ensure your project and publish profiles are committed to your version control repository, so they are accessible to your CI/CD pipeline.

Manage Variables:

Use your CI/CD tool's capabilities to manage deployment variables specific to each environment. These variables could include connection strings, credentials, or other configuration settings. Map these variables to the ones defined in your publish profiles.

Select Publish Profile:

In your CI/CD pipeline, specify which publish profile to use for deployment based on the target environment. This ensures that the appropriate variables and configurations are applied during deployment.

4) Create the build pipeline using git, solution file and publish profile.

4.1) Create the build pipeline as it is as per the brief description given at the beginning.

This step is to take the solution file that you push from your SQL project from visual studio code to GitHub and create an artifact for the release pipeline to work with.

5) Testing Integration:

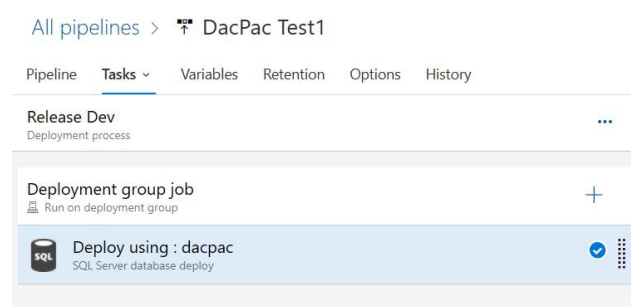
For testing you can add files to your GIT and see if that triggers the build pipeline, if it does, we can check for the artifact drop. If the build pipeline runs seamlessly then we will have an artifact drop with the new solution file. Using this we aim to ensure the reliability and stability of the software.

Refer to the brief description at the beginning for screenshots.

6) Create Release pipeline and add publish profile location for different environments.

6.1) **Deployment Setups:** We will set up deployment/release pipelines, defining deployment stages and environment configurations. Azure DevOps provides integration with Azure services, allowing us to deploy applications to Azure App Services, Kubernetes clusters, or other relevant platforms.

Jobs required for the agent: Deploy using DACPAC



The Task we used has the following properties:

SQL Server database deploy ⓘ
[View YAML](#)
[Remove](#)

Task version 0.* ▼

Display name *

Deploy using : dacpac

Deploy SQL Using * ⓘ

Sql Dacpac ▼

DACPAC File * ⓘ

\$(System.DefaultWorkingDirectory)/_EDM Impact-
CI/drop/AdventureWorks2019.dacpac

...

Specify SQL Using * ⓘ

Connection String ▼

Connection String * ⓘ

\$(ConnectionString)

Additional Arguments ⓘ

/p:DropObjectsNotInSource=True /p:BlockOnPossibleDataLoss=False

Control Options ▼

Output Variables ▼

6.2) Inputting Publish Profile according to the environment (Dev, UAT, Prod)

In the above given sample pipeline properties, you can see the publish profile argument highlighted. Put the publish profile location from your artifact drops here to define the variables in your solution file.

This will vary from environment to environment as per the DB names such as one variable in solution as @Warehouse can refer to multiple DBs when deployed such as Warehouse_Dev, Warehouse_UAT, Warehouse_Prod depending on the requirement. This is controlled via publish profiles.

Conclusion:

This Proof of Concept aims to demonstrate the viability of creating an instructional manual that effectively guides users through complex processes or systems. By focusing on clarity, user-centric design, and feedback-driven iteration, this POC will lay the foundation for developing comprehensive instructional materials that empower users and enhance their experiences.